

APAN PS5400: Managing Data

Week 3: Creating Relational Databases

Lecturer: François Scharffe

Recap of last week

- Relational data model
- Relational databases
- Relational Database Management Systems (RDBMS)
- SQL as a query language

This week

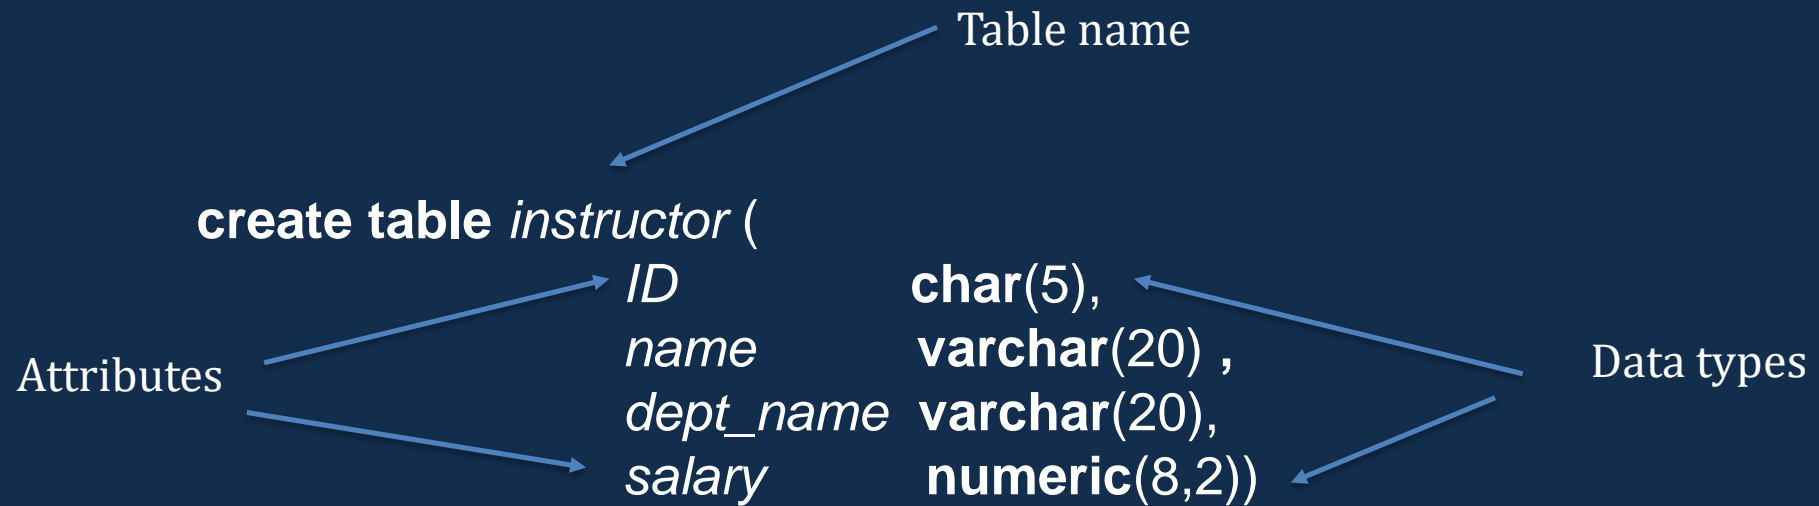
- Creating relational databases using SQL
- Primary key of a table
- Foreign key constraints and referential integrity
- Other types of constraints
- Insertion of data into tables
- Modifying data
- Deleting tuples
- Using Postgres SQL to create databases on your computer

SQL as Data Definition Language

Can be used to create relations as well as additional information about relations, such as:

- The schema of the relation
- The data types of the values to be stored in the relation
- The integrity constraints on the relation
- The set of indices on the relation
- Information about access privileges and security
- The physical storage structure of the relation on the disk

Create table



Most basic form. We can add various types of constraints to this basic form. See next few slides.

Create table with primary key constraint

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) ,  
    dept_name   varchar(20),  
    salary      numeric(8,2),  
    primary key (ID),  
    )
```

← Primary key constraint

What is a primary key? See next.

Primary keys

- Recall that a table (i.e., a relation) is just a set of tuples and a set cannot have duplicate tuples.
- This constraint on tables can be enforced by making a set, K , of attributes of the relation, R , such that no two tuples in R can have the same set of values on K .
 - In the *create table* statement of the previous slide we specify that no two tuples of the *instructor* table can have the same value on ID attribute.

Primary Keys: Formal definition

- Let $K \subseteq R$, where R is the attributes of a relation
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple in each possible instance of R .
 - Which implies that there cannot be two tuples in any instance of R with the same values on K , given that a set cannot have duplicates.
 - E.g., $\{ID\}$ and $\{ID, name\}$ are both superkeys of R
- A superkey K is a **candidate key** if no proper subset of it is itself a superkey (i.e., K is minimal)
- Any candidate key can be chosen as the **primary key** for R
 - ID is a candidate key and, thus, can be the primary key
 - If there are more than one candidate keys, then choose the one on which you want to build an index as the primary key

What if we don't declare a primary key when creating a table? Then will the table allow duplicate rows?

How do you decide what should be primary key? (Recall we can have a PK that is a composite of several attributes.) It depends on the rules of the organization (or the world) that the database is modelling?

Enrollment: (Student, Course, Section, Semester)

- What if you make Student the PK?
- What if you make {Student, Course} the PK?
- What if you make {Student, Course, Section} the PK?
- What if you make {Student, Course, Semester} the PK?
- What if you make {Student, Semester} the PK?
- What if you make {Student, Course, Section, Semester} the PK?

Create Table with Foreign Key

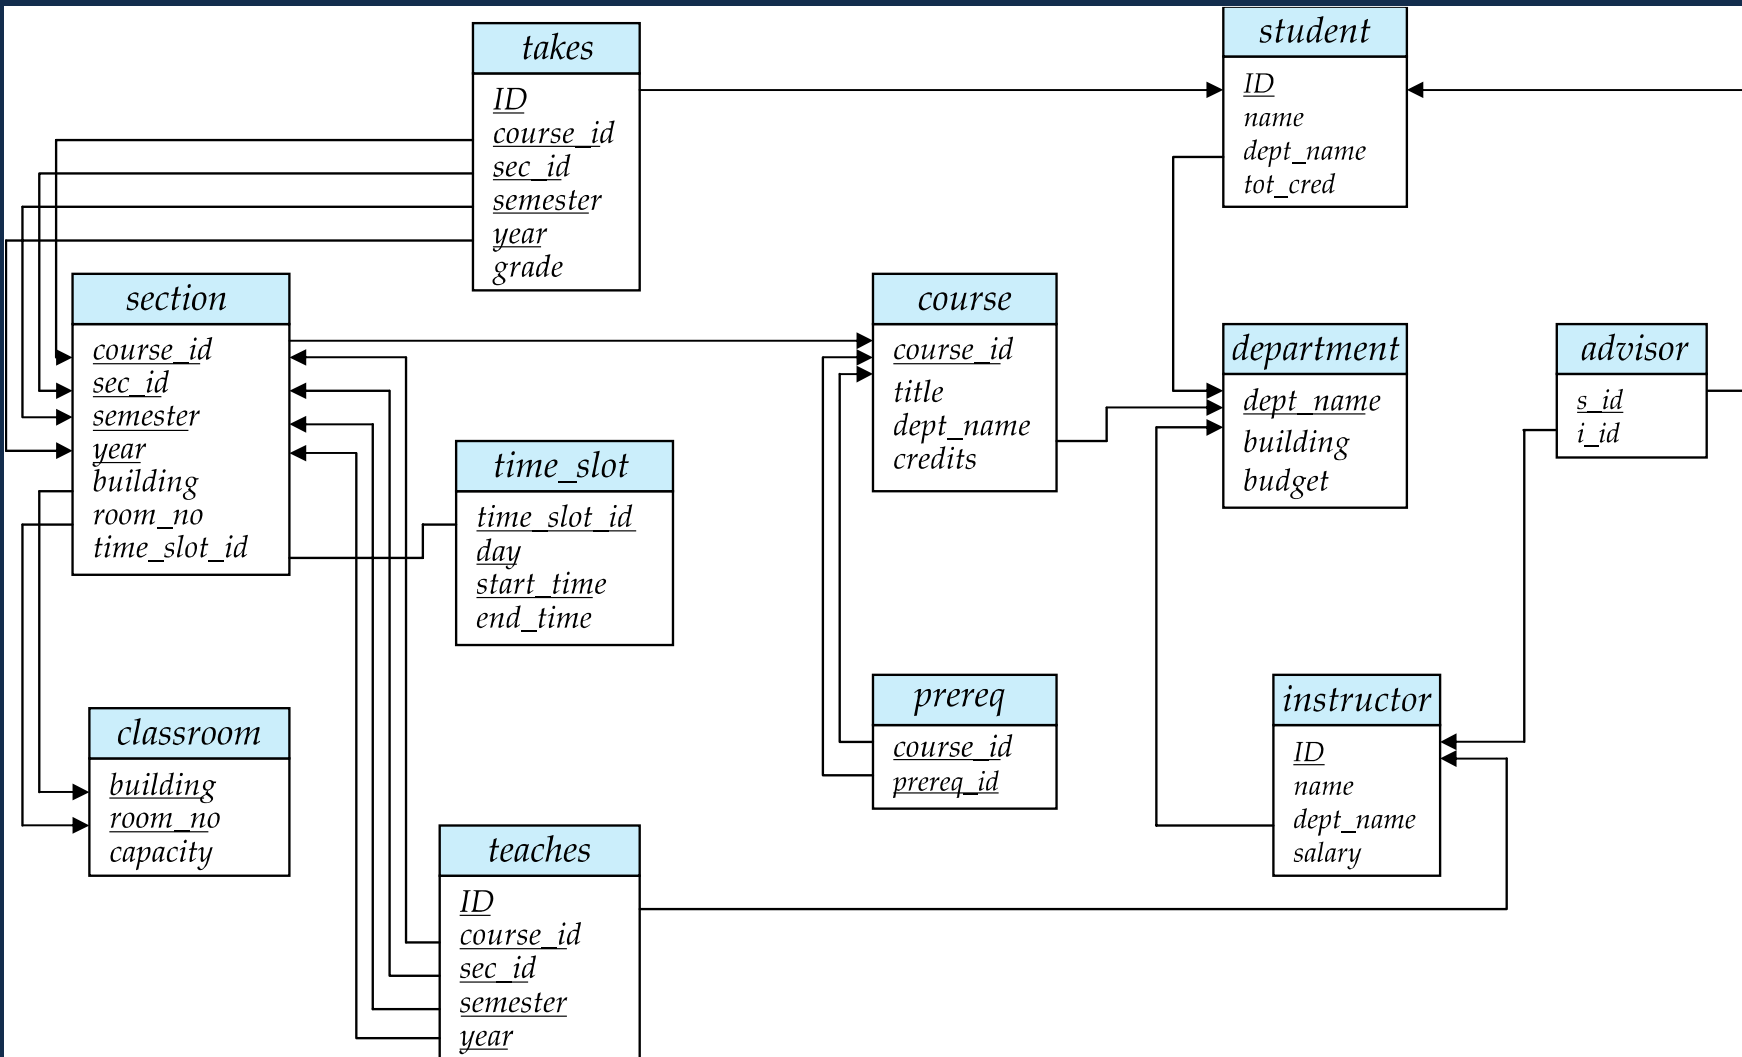
```
create table instructor (  
    ID          char(5),  
    name       varchar(20) not null,  
    dept_name varchar(20),  
    salary    numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references department)
```

Foreign key constraint



Foreign key constraint and Referential Integrity

- A table (e.g., *instructor*) may have an attribute (e.g., *department_name*) such that whether a value of that attribute in that table is a valid value or not has to be checked in terms of whether that value occurs in another table (e.g., *department*)
- The other table is regarded as the validator of such values in any other table.
- In this case we say there is a foreign key reference from any table which has that attribute to the validator table.
- A database in which all foreign key constraints are respected is said to have **referential integrity**.
- In the SQL code for creating a table we can specify what action to take if insertion of data violates a foreign key constraint on table: Block the insertion? Take a corrective action on the validator table? No action?
 - This will be covered in the elective course on SQL.



The schema of an RDB with foreign key constraints. An arrow from a table to another table indicates that there is a foreign key constraint from the first table to the second table.

Thus, we see that there is a foreign key reference from *course* to *department*.

Create table with not null constraint

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) not null, ← Not null constraint  
    dept_name   varchar(20),  
    salary      numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references department)
```

This constraint specifies that *name* cannot take null values.
Primary key field is not allowed to take null values.

Null values

- Some of the attributes in some of the tuples in a relation can have **null values**.
- null denotes an unknown value or a value that doesn't exist.
- Arithmetic expressions with a null value results in a null value
- Aggregate functions simply ignore null values
- P is not null return false if P is null, and true otherwise

Other constraints

- Check constraints
- Assertions
- Triggers
- Not all RDBMS support these constraints
- These will be covered in the SQL elective course

Inserting data,

- insert into course values ('APAN5400', 'Managing Data', 'APAN', 3), or equivalently
- insert into course (*course_id, title, dept_name, credits*) values ('APAN5400', 'Managing Data', 'APAN', 3)
- Add a new tuple to student with tot_cred set to null
insert into *student* values (' 3003' , ' Green' , ' Finance' , null);

Updates

The general form is:

```
update table_name  
set attribute_value = value  
where boolean condition
```

Example:

```
update instructor  
set salary = salary * 1.5  
where salary < 50,000
```

Deletes

- General form:
delete from *table* where *boolean condition*
- Example:
delete from *instructor* where *salary > 75000*

Demo

- We will spend the rest of the class doing a demo of how to install Postgres on a computer, and
- Creating a sample database
- Pay attention—the first part of your first project (Project 1.1) requires you to master this.

Recap: What we covered this week

- SQL to create tables
- Primary key of a table
- Foreign key constraints and referential integrity
- Other types of constraints
- Insertion, update, delete

Next week

- Views
- Queries using views
- Functional dependencies
- BCNF
- Decomposition of tables to make database BCNF compliant